

Big Data Computing Final Project Guidelines

Gabriele Tolomei

Sapienza University of Rome, Italy

Email: tolomei@di.uniroma1.it Homepage: <https://www.di.uniroma1.it/~tolomei>

Scope of the Document

This document provides students of the **Big Data Computing** class (hereinafter, “you”) with a list of *guidelines* for developing their projects, which are mandatory for final grading.

1 Project Requirements

Projects must of course refer to a typical big data task, such as those seen during classes: e.g., clustering, regression/classification, recommendation, graph analysis, using large datasets in *any* application domain of interest.

The development environment must be the same of that used throughout the course, namely PySpark in combination with Google Colab.

Projects can be done either individually or in group of **at most** 2 students, and they should be accompanied by a brief presentation written in english (e.g., a few PowerPoint slides).

There is no restriction on the learning technique(s) you should use to solve your task (e.g., K-means, logistic regression, matrix factorization, artificial neural networks, etc.) In other words, you can experiment with any technique you want – providing it is relevant to your goal and hopefully comparing more than one with each other – including those not covered in class.

2 Project Selection

There exist several resources which contain many project ideas as well as the related datasets to work on; among those, I would recommend the following ones:

- [Kaggle](#)
- [UCI Machine Learning Repository](#)
- [Awesome Public Datasets](#)

Alternatively, you can also come up with your own project idea, as long as it satisfies the requirements indicated in the section above. Eventually, *all* the project proposals must anyway be agreed with the professor in advance.

3 Project Evaluation


Projects will be evaluated according to the following **5** criteria:

- **“Big-Data-ness”**: This criterion will check if the “big data” requirements of the project are actually met. It will provide answers to questions like “*How complex the task is?*” or “*How large is the dataset used?*”.

- **Coding:** This criterion will evaluate the style of programming. It will provide answers to questions like “*Is the source code easy to read, reusable, and well commented?*” or “*How much effort would it take for the source code to get into a hypothetical production system?*”.
- **Methodology:** This criterion will assess the soundness of the methodological approach used to achieve the project’s goal. It will provide answers to questions like “*Does the pipeline implement all the stages needed?*” or “*How reliable are the results obtained?*”.
- **Originality/Impact:** This criterion will validate both the novelty of the task that the project aims to solve and the creativity of the solutions implemented. It will provide answers to questions like “*Is the task novel/impactful or has it been already deeply investigated?*” or “*How conventional are the solutions adopted to solve the task?*”.
- **Presentation:** This criterion will judge the quality of the presentation which is attached to the project. It will provide answers to questions like “*Is the project well defined and its goal clearly stated?*” or “*What is the level of understanding of the challenges associated with the specific task?*”.

4 Project Submission

Projects will be submitted for grading to the [Moodle web page](#) of the course. More specifically, for each examination session there will be an *Assignment* entry on Moodle, which will allow you to upload your project material. Moodle assignments are identified by the following icon:



For example, suppose you are ready to submit your project on the June session. On the Moodle web page of the course, you will see an assignment corresponding to that session (e.g., called “*June Session Project Submissions*”). This assignment will have an internal deadline, e.g., **June, 20 at 23:59:59**. This means that everyone who wants to submit her/his project on the June session must upload the material within the specified deadline.

The project material should preferably be packaged inside a (compressed) folder (e.g., **.tar**, **.tgz**, **.bz2**, etc.) containing the following items:

- A Google Colab notebook with all the source code of the project;
- A presentation (e.g., PowerPoint or PDF slides) with a description of the project, the main choices made to accomplish the task, and the results obtained.

Please, consider that the notebook must be “ready-to-execute”: in other words, it must contain everything to be run properly and successfully (e.g., environment setup, library dependencies, dataset collection, etc.)

To ease the grading process, it may be helpful to setup a naming convention for the project submission. The uploaded folder should be named as follows: **X(_Y).Z**, where:

- X and Y are *student IDs* of the project team;
- Z = Archive extension (e.g., **tar**, **tgz**, **bz2**, etc.).

In the case of a single-person team, the folder will be named as **X.Z**, whilst in the case of two-person team, X and Y will be the ID of the first and second student, *after they have been alphabetically ordered by their last name*.

For example, if the project is done by a single student whose ID is “12345”, then $X=12345$ and the folder will be named `12345.tgz` (or similar extension). Instead, if the project is done by two students: *Clark Kent* whose student ID is “67890” and *Bruce Wayne* whose student ID is “12345”, then the project folder will be named `67890_12345.tgz` (or similar extension). In addition, files within the project folder (i.e., the Google Colab notebook and the presentation) should follow the same naming convention.

5 Project Discussion

Right after the project submission deadline of an examination session and before the next one, it will be setup an oral discussion slot for each of the submitted projects.

The main discussion will be in english: you will be asked to show a working demo of your project by running (part of) your notebook, and to motivate the choices you made. Questions about any other topic addressed during the course may also be asked, but those can be answered either in english or in italian, as you prefer.

[Due to the current restrictions posed by the COVID-19 emergency, it is quite hard to predict now whether oral discussions will be done face to face or remotely. Either way, you will be notified in advance on how to attend the examination.]

6 Honor Code

The basic principle under which you are expected to operate is that you should submit your own work. Of course, if you are part of a 2-member project team this principle will naturally extend to the team as a whole. More specifically, attempting to take credit for someone else's work by turning it in as your own constitutes *plagiarism*, which is a serious violation of basic academic standards.

To observe the honor code, you are invited to follow the rules below:

- You must not submit or look at solutions or program code that are not your own;
- You must not share your solution code with other students, nor ask others to share their solutions with you;
- You must indicate on your submission any assistance you received;

Please be aware that all submissions are subject to automated plagiarism detection.

As a final remark, many forms of collaboration - if legitimate and properly acknowledged - are acceptable and indeed encouraged.