

# Big Data Computing

Master's Degree in Computer Science

2019-2020

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)



SAPIENZA  
UNIVERSITÀ DI ROMA

# Recap from Last Lecture

- MapReduce → new distributed computing framework suitable for working with large scale datasets
- Useful in all those situations where data need to be accessed sequentially
- May be hard to program and does not support well multiple map-reduce rounds

# Data-Flow Systems

- MapReduce uses 2 "**ranks**" of tasks: one for map the other for reduce

# Data-Flow Systems

- MapReduce uses 2 "**ranks**" of tasks: one for map the other for reduce
- Data flows from the first rank to the second

# Data-Flow Systems

- MapReduce uses 2 "**ranks**" of tasks: one for map the other for reduce
- Data flows from the first rank to the second
- Generalized Data-Flow Systems abstract from this in two ways:
  - Allow any number of "ranks"/tasks
  - Allow functions other than just map and reduce

# Data-Flow Systems

- MapReduce uses 2 "**ranks**" of tasks: one for map the other for reduce
- Data flows from the first rank to the second
- Generalized Data-Flow Systems abstract from this in two ways:
  - Allow any number of "ranks"/tasks
  - Allow functions other than just map and reduce
- As long as data goes in one direction only, recovery at intermediate rank is possible

# Spark: Most Popular Data-Flow System

- Expressive computing framework not limited to map-reduce model

# Spark: Most Popular Data-Flow System

- Expressive computing framework not limited to map-reduce model
- In addition to MapReduce, Spark provides:
  - Fast data sharing (no intermediate saving to local disks + caching)
  - General execution graphs (DAGs)
  - Richer functions than just map and reduce



# Spark: Most Popular Data-Flow System

- Expressive computing framework not limited to map-reduce model
- In addition to MapReduce, Spark provides:
  - Fast data sharing (no intermediate saving to local disks + caching)
  - General execution graphs (DAGs)
  - Richer functions than just map and reduce
- Compatible with Hadoop

# Spark: Introduction

- Originally developed at UC Berkeley in 2009 and later donated to the Apache Software Foundation (open-source)

# Spark: Introduction

- Originally developed at UC Berkeley in 2009 and later donated to the Apache Software Foundation (open-source)
- Implemented in **Scala** (running on top of the Java Virtual Machine)

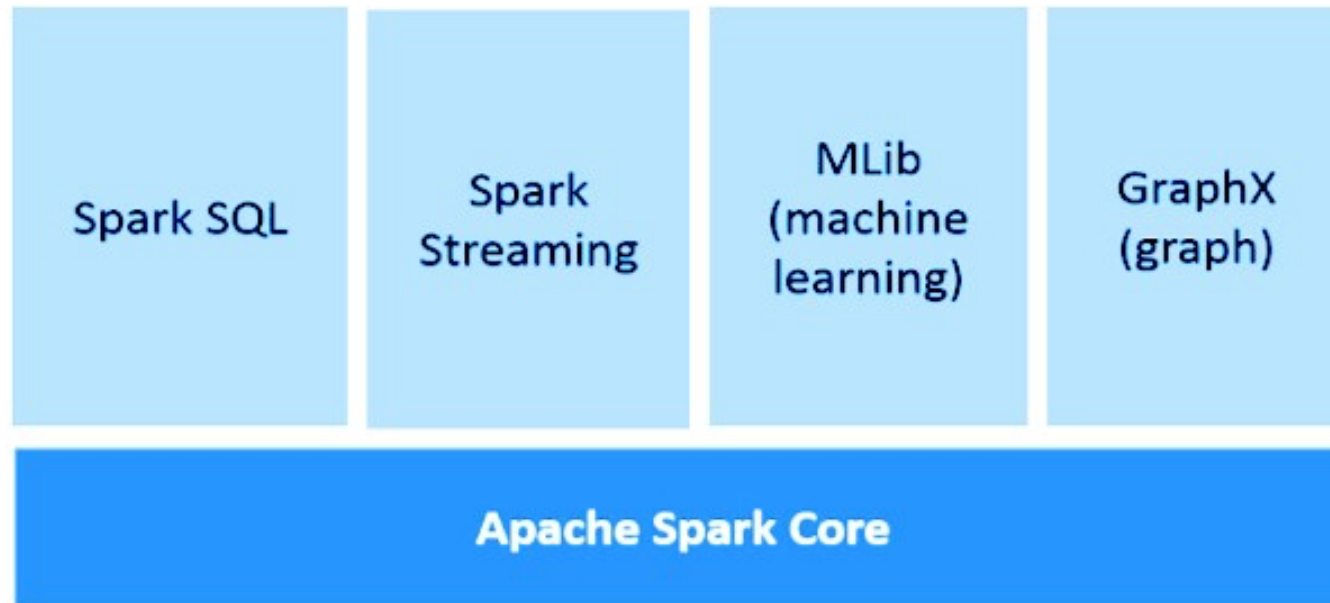
# Spark: Introduction

- Originally developed at UC Berkeley in 2009 and later donated to the Apache Software Foundation (open-source)
- Implemented in **Scala** (running on top of the Java Virtual Machine)
- Unified **computing engine** (**Spark Core**)

# Spark: Introduction

- Originally developed at UC Berkeley in 2009 and later donated to the Apache Software Foundation (open-source)
- Implemented in **Scala** (running on top of the Java Virtual Machine)
- Unified **computing engine** (**Spark Core**)
- Set of **high-level APIs** for data analysis:
  - **Spark SQL** (structured data), **MLib** (machine learning), **GraphX** (graph analytics), **Spark Streaming** (stream data processing)

# Spark: Overview



# Spark: Introduction

- Unlike Hadoop, Spark does not come with a storage system

# Spark: Introduction

- Unlike Hadoop, Spark does not come with a storage system
- In fact, it provides interfaces for many local and distributed storage systems:
  - HDFS, Amazon S3, Cassandra, Hive Metastore, or classical RDBMS



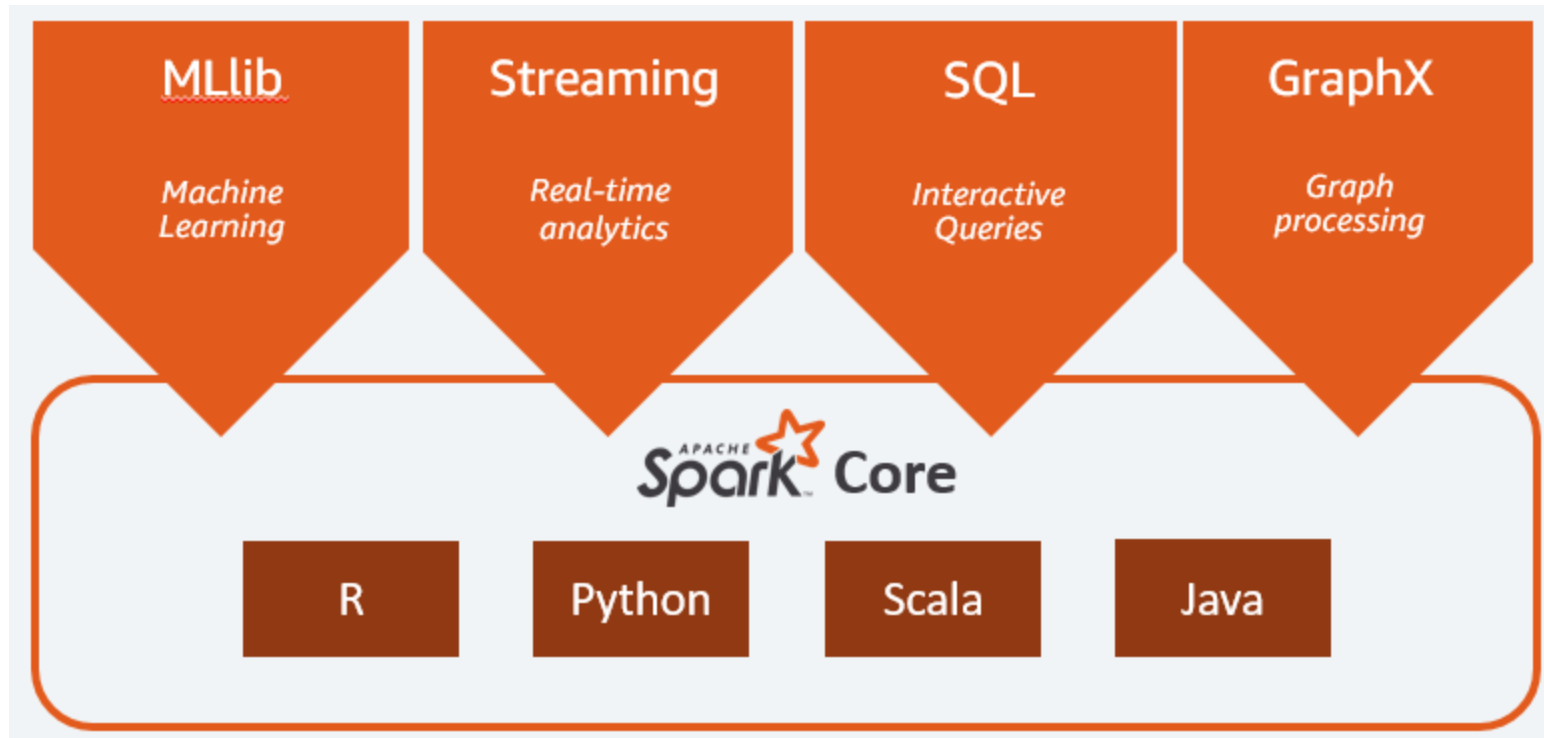
# Spark: Introduction

- Unlike Hadoop, Spark does not come with a storage system
- In fact, it provides interfaces for many local and distributed storage systems:
  - HDFS, Amazon S3, Cassandra, Hive Metastore, or classical RDBMS
- Additionally, Spark's APIs are available for many programming languages: Scala, Java, Python, and R

# Spark: Introduction

- Unlike Hadoop, Spark does not come with a storage system
- In fact, it provides interfaces for many local and distributed storage systems:
  - HDFS, Amazon S3, Cassandra, Hive Metastore, or classical RDBMS
- Additionally, Spark's APIs are available for many programming languages: Scala, Java, Python, and R
- This flexibility is the key of its success in the Big Data landscape

# Spark: More Detailed Overview



# Spark: Features

- Fault-tolerant system

# Spark: Features

- Fault-tolerant system
- In-memory caching which enables efficient execution of multi-round algorithms (i.e., multiple sequential tasks)
  - performance improvement w.r.t. Hadoop

# Spark: Features

- Fault-tolerant system
- In-memory caching which enables efficient execution of multi-round algorithms (i.e., multiple sequential tasks)
  - performance improvement w.r.t. Hadoop
- Spark can run:
  - on a single machine → local mode
  - on a cluster managed by a cluster manager (e.g., Spark Standalone, YARN, Mesos)

# Spark: Features

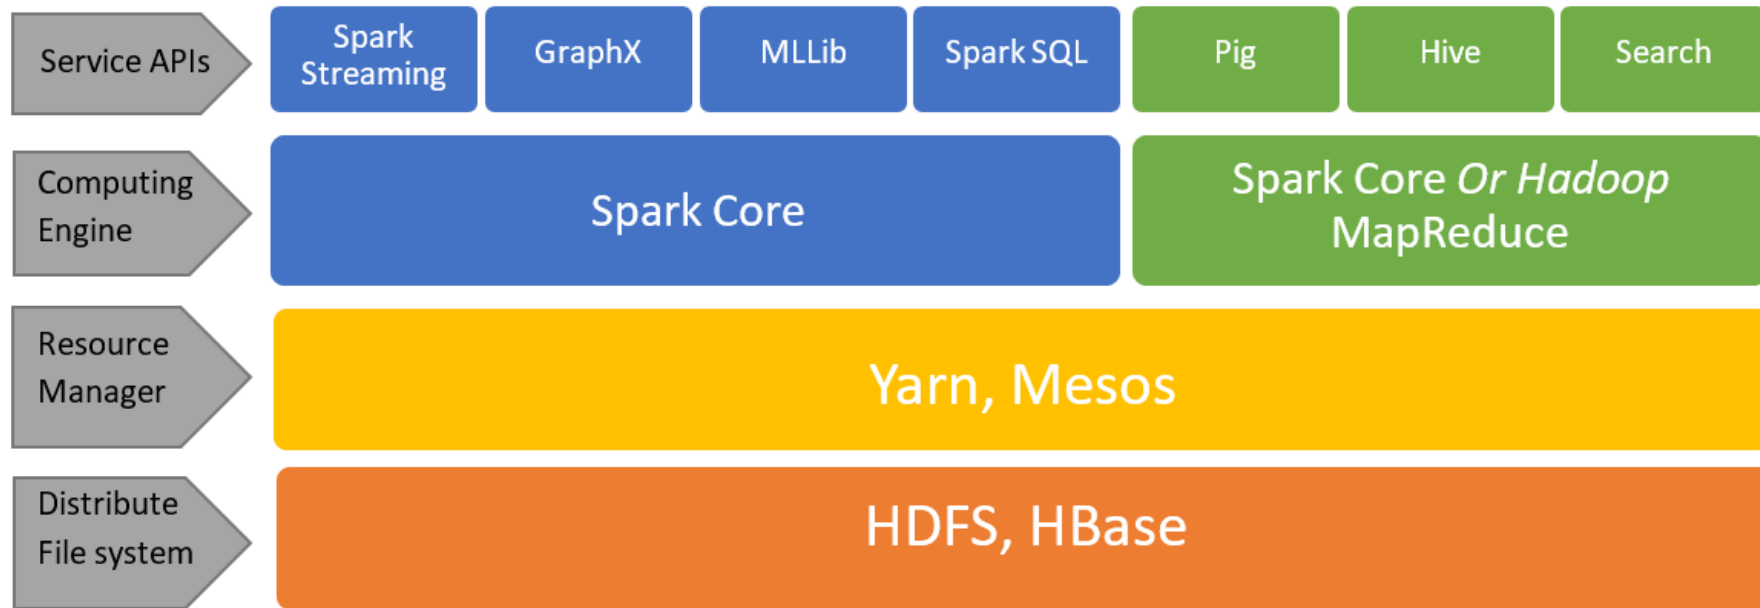


Figure 1 - Spark Context

# Spark Application: Driver

- The **driver process** (a.k.a. **master** in MapReduce terminology) runs the application's entry point from a node in the cluster



# Spark Application: Driver

- The **driver process** (a.k.a. **master** in MapReduce terminology) runs the application's entry point from a node in the cluster
- The driver is responsible for:
  - Maintaining information about the application
  - Responding to a user program or input
  - Analyzing, distributing, and scheduling work across executors

# Spark Application: Driver

- The **driver process** (a.k.a. **master** in MapReduce terminology) runs the application's entry point from a node in the cluster
- The driver is responsible for:
  - Maintaining information about the application
  - Responding to a user program or input
  - Analyzing, distributing, and scheduling work across executors
- The driver is represented by an object called **Spark Context**

# Spark Application: Executor(s) and Cluster Manager

- **Executor processes** (a.k.a. **workers** in Hadoop terminology) actually compute the tasks assigned by the driver


# Spark Application: Executor(s) and Cluster Manager

- **Executor processes** (a.k.a. **workers** in Hadoop terminology) actually compute the tasks assigned by the driver
- Each executor is responsible for:
  - Running the code assigned to it by the driver
  - Reporting the state of the computation back to the driver

# Spark Application: Executor(s) and Cluster Manager

- **Executor processes** (a.k.a. **workers** in Hadoop terminology) actually compute the tasks assigned by the driver
- Each executor is responsible for:
  - Running the code assigned to it by the driver
  - Reporting the state of the computation back to the driver
- The cluster manager controls physical machines and allocates resources to applications

# Spark Application



# Spark Application: Considerations

- Driver and executors are processes which can live on the same machines or on different nodes

# Spark Application: Considerations

- Driver and executors are processes which can live on the same machines or on different nodes
- When Spark is running in local mode, both the driver and executors are running as separate threads on the same machine



# Spark Application: Considerations

- Driver and executors are processes which can live on the same machines or on different nodes
- When Spark is running in local mode, both the driver and executors are running as separate threads on the same machine
- Executors mostly run Scala code

# Spark Application: Considerations

- Driver and executors are processes which can live on the same machines or on different nodes
- When Spark is running in local mode, both the driver and executors are running as separate threads on the same machine
- Executors mostly run Scala code
- Driver can be governed by different languages using Spark's APIs

# Resilient Distributed Dataset (RDD)

- Fundamental **abstraction** of Spark to indicate a collection of elements of the same type
  - Generalization of MapReduce's key-value pairs

# Resilient Distributed Dataset (RDD)

- Fundamental **abstraction** of Spark to indicate a collection of elements of the same type
  - Generalization of MapReduce's key-value pairs
- RDDs are **partitioned** and possibly spread across multiple nodes of the cluster

# Resilient Distributed Dataset (RDD)

- Fundamental **abstraction** of Spark to indicate a collection of elements of the same type
  - Generalization of MapReduce's key-value pairs
- RDDs are **partitioned** and possibly spread across multiple nodes of the cluster
- Best suited for applications that apply the same operation across all the elements of the dataset

# RDD: Partitions

- Each RDD is split into chunks called **partitions** distributed across nodes

# RDD: Partitions

- Each RDD is split into chunks called **partitions** distributed across nodes
- A program can specify the number of partitions for an RDD (otherwise Spark will choose one)

# RDD: Partitions

- Each RDD is split into chunks called **partitions** distributed across nodes
- A program can specify the number of partitions for an RDD (otherwise Spark will choose one)
- Programmer can also decide whether to use the default **Hash Partitioner** or a custom one



# RDD: Partitions

- Each RDD is split into chunks called **partitions** distributed across nodes
- A program can specify the number of partitions for an RDD (otherwise Spark will choose one)
- Programmer can also decide whether to use the default **Hash Partitioner** or a custom one
- A typical number of partitions is 2 or 3 times the number of cores

# RDD: Partitions

- Partitioning enables the following:
  - **Data reuse** → data is kept in executors' main memory so as to avoid expensive access to external disks

# RDD: Partitions

- Partitioning enables the following:
  - **Data reuse** → data is kept in executors' main memory so as to avoid expensive access to external disks
  - **Parallelism** → Some data transformations are applied independently to each partition thereby avoiding expensive data transfers

# RDD: Characteristics

- RDDs are **immutable** (i.e., read-only)

# RDD: Characteristics

- RDDs are **immutable** (i.e., read-only)
- Can be created either from data stored on distributed file system (e.g., HDFS) or as a result of transformations of other RDDs

# RDD: Characteristics

- RDDs are **immutable** (i.e., read-only)
- Can be created either from data stored on distributed file system (e.g., HDFS) or as a result of transformations of other RDDs
- RDDs do not need to be always materialized
  - Each RDD maintains a sort of "trace" of transformations (lineage) that led to the current status
  - This way, RDD can always be re-created even upon a failure

# RDD Operations

- Let  $A$  be an RDD, the following **3 operations** are possible:

# RDD Operations

- Let A be an RDD, the following **3 operations** are possible:
  - **Transformations** → generate a new RDD B from the data in A



# RDD Operations

- Let A be an RDD, the following **3 operations** are possible:
  - **Transformations** → generate a new RDD B from the data in A
  - **Actions** → launch a computation on the data in A, which returns a value to the application

# RDD Operations

- Let A be an RDD, the following **3 operations** are possible:
  - **Transformations** → generate a new RDD B from the data in A
  - **Actions** → launch a computation on the data in A, which returns a value to the application
  - **Persistence** → save the RDD in memory for later actions

# RDD Operations: Transformations

- **Narrow:** each partition of A contributes at most to one partition of B (e.g., **map**)
  - No need to shuffle data across nodes

# RDD Operations: Transformations

- **Narrow:** each partition of A contributes at most to one partition of B  
(e.g., **map**)
  - No need to shuffle data across nodes
- **Wide:** each partition of A may contribute to multiple partitions of B  
(e.g., **groupBy**)
  - Possible need to transfer data across nodes


# RDD Operations: Transformations

- **Narrow:** each partition of A contributes at most to one partition of B  
(e.g., **map**)
  - No need to shuffle data across nodes
- **Wide:** each partition of A may contribute to multiple partitions of B  
(e.g., **groupBy**)
  - Possible need to transfer data across nodes
- **Lazy evaluation:** nothing is computed unless required by an action

# Narrow vs. Wide Transformations

Narrow

Input and output stay  
on the same partition




# Narrow vs. Wide Transformations


Narrow

Input and output stay  
on the same partition

No data transfers




# Narrow vs. Wide Transformations



Input from other partitions may be needed



# Narrow vs. Wide Transformations




Input from other partitions may be needed

Data shuffling across nodes

# Narrow vs. Wide Transformations


Narrow



Input and output stay  
on the same partition

No data transfers

Wide



Input from other  
partitions may be needed

Data shuffling  
across nodes

# RDD Operations: Actions

- Example: the count method returns the number of elements of the RDD
- When the action is called the RDD is actually materialized (lazy evaluation)

# Spark DataFrame and Dataset APIs

- RDDs are the most basic data model used by Spark
  - low-level and schema-less

# Spark DataFrame and Dataset APIs

- RDDs are the most basic data model used by Spark
  - low-level and schema-less
- On top of RDD API, **Spark SQL** module provides **2 interfaces** to operate on structured data like tables in relational databases:
  - **DataFrame API**
  - **Dataset API**

# Spark: DataFrame

- Distributed collection of data organized into **named columns**

# Spark: DataFrame

- Distributed collection of data organized into **named columns**
- Allows higher level abstraction than plain vanilla RDDs

# Spark: DataFrame

- Distributed collection of data organized into **named columns**
- Allows higher level abstraction than plain vanilla RDDs
- Since Spark 2.0 it is part of a more general Dataset API
  - Dataset API is available only for Scala and Java as it extends DataFrame API with type-safe, object-oriented programming interface



# Spark: DataFrame

- Distributed collection of data organized into **named columns**
- Allows higher level abstraction than plain vanilla RDDs
- Since Spark 2.0 it is part of a more general Dataset API
  - Dataset API is available only for Scala and Java as it extends DataFrame API with type-safe, object-oriented programming interface
- Similar to [Pandas DataFrame](#) unless few differences

# Spark DataFrame vs. Pandas DataFrame

- Spark DataFrames are **immutable**: once created they cannot be modified

# Spark DataFrame vs. Pandas DataFrame

- Spark DataFrames are **immutable**: once created they cannot be modified
- As for RDDs, Spark may apply 2 kinds of operations on DataFrames:  
**transformations** and **actions**

# Spark DataFrame vs. Pandas DataFrame

- Spark DataFrames are **immutable**: once created they cannot be modified
- As for RDDs, Spark may apply 2 kinds of operations on DataFrames: **transformations** and **actions**
- Lazy evaluation allows to queue transformations applied to elements of a DataFrame until an action is called

# Spark DataFrame vs. Pandas DataFrame

- Spark DataFrames are **immutable**: once created they cannot be modified
- As for RDDs, Spark may apply 2 kinds of operations on DataFrames: **transformations** and **actions**
- Lazy evaluation allows to queue transformations applied to elements of a DataFrame until an action is called
- DataFrame (and Dataset as well) can be turned back to RDD

# Spark vs. Hadoop MapReduce

- **Performance:** Spark is usually faster
  - In-memory data processing vs. data persistencing to disk after any map/reduce step
  - Spark requires lots of memory to run fast, otherwise its performance deteriorates
  - MapReduce integrates better with other services running

# Spark vs. Hadoop MapReduce

- **Performance:** Spark is usually faster
  - In-memory data processing vs. data persistencing to disk after any map/reduce step
  - Spark requires lots of memory to run fast, otherwise its performance deteriorates
  - MapReduce integrates better with other services running
- **Ease of use:** Spark provides a higher-level API which is easier to program

# Spark vs. Hadoop MapReduce

- **Performance:** Spark is usually faster
  - In-memory data processing vs. data persistencing to disk after any map/reduce step
  - Spark requires lots of memory to run fast, otherwise its performance deteriorates
  - MapReduce integrates better with other services running
- **Ease of use:** Spark provides a higher-level API which is easier to program
- **Data processing:** Spark is more flexible and general



# Take-Home Message of Today

- Spark is a general-purpose distributed data processing engine which overcomes many of the Hadoop's limitations

# Take-Home Message of Today

- Spark is a general-purpose distributed data processing engine which overcomes many of the Hadoop's limitations
- Spark provides a rich ecosystem of services to work on (big) data through APIs accessible via multiple programming languages

# Take-Home Message of Today

- Spark is a general-purpose distributed data processing engine which overcomes many of the Hadoop's limitations
- Spark provides a rich ecosystem of services to work on (big) data through APIs accessible via multiple programming languages
- Spark's **DataFrame** as the main abstraction for playing with data