

Big Data Computing

Master's Degree in Computer Science

2019-2020

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

tolomei@di.uniroma1.it



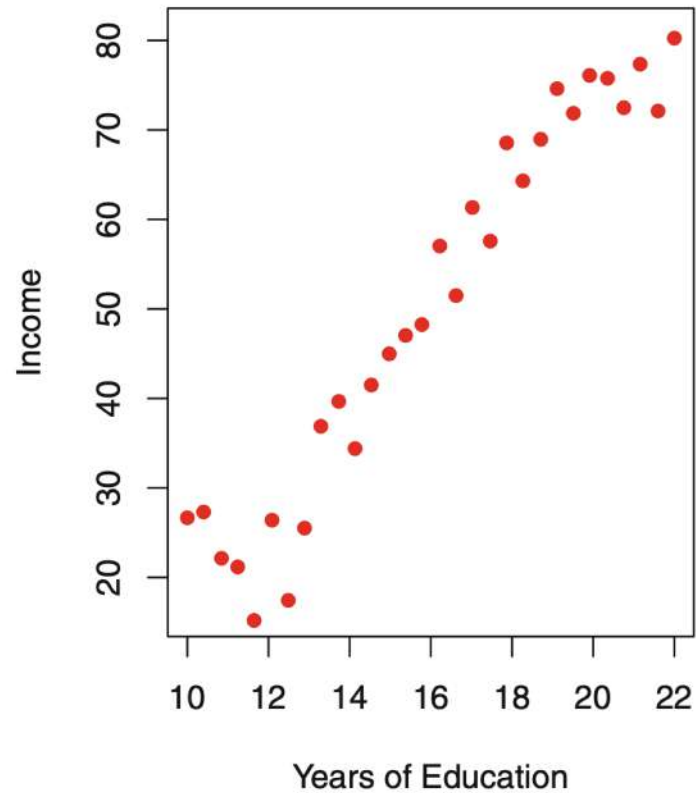
SAPIENZA
UNIVERSITÀ DI ROMA

Recap from Last Lecture

- Supervised Learning as an optimization problem
 - Hypothesis space (assumption)
 - Loss Function (objective)
 - Learning Algorithm (optimizer)
- Regression vs. Classification
- Bias-Variance Tradeoff
- Model selection vs. Model evaluation

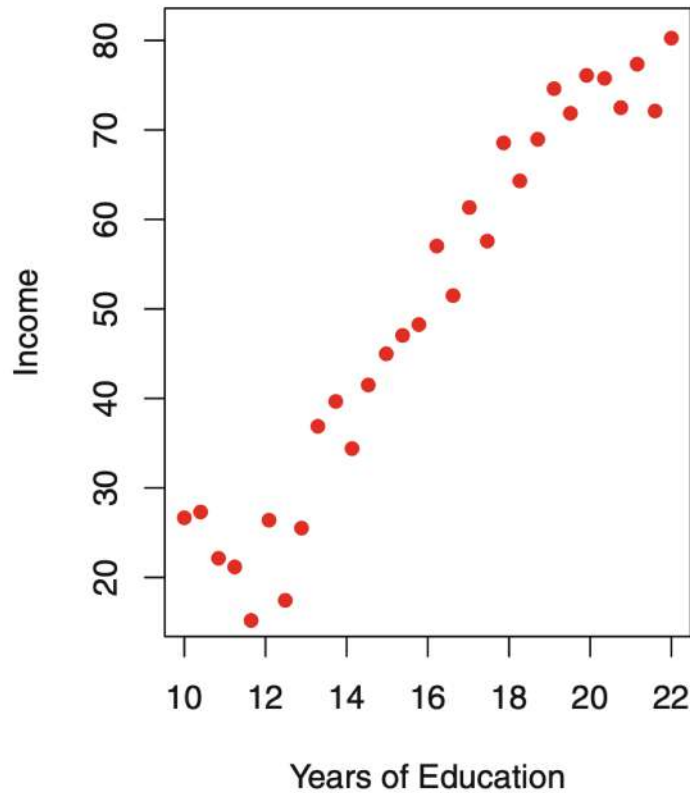
LINEAR REGRESSION

Example: Y =Income vs. X =Education

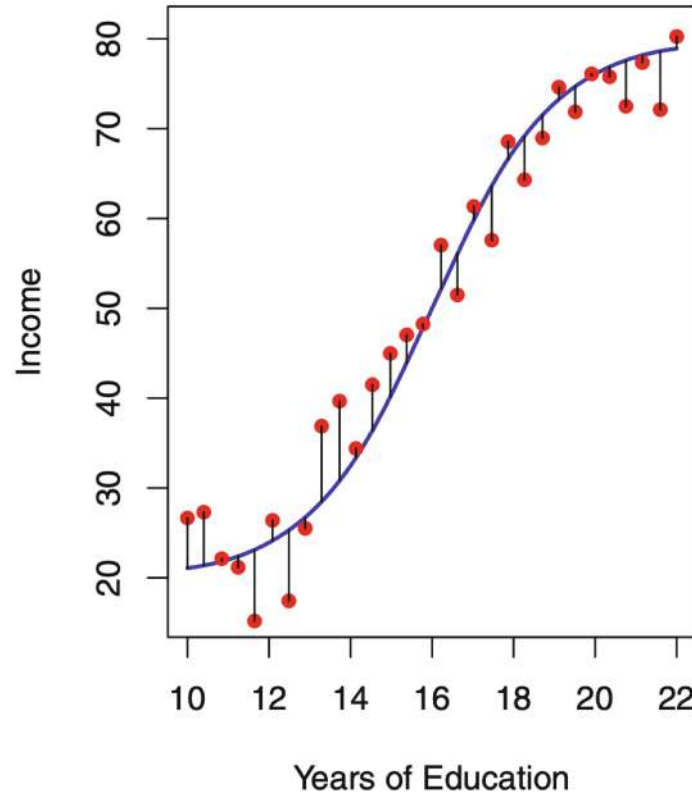


Observations
(simulated)

Example: Y =Income vs. X =Education



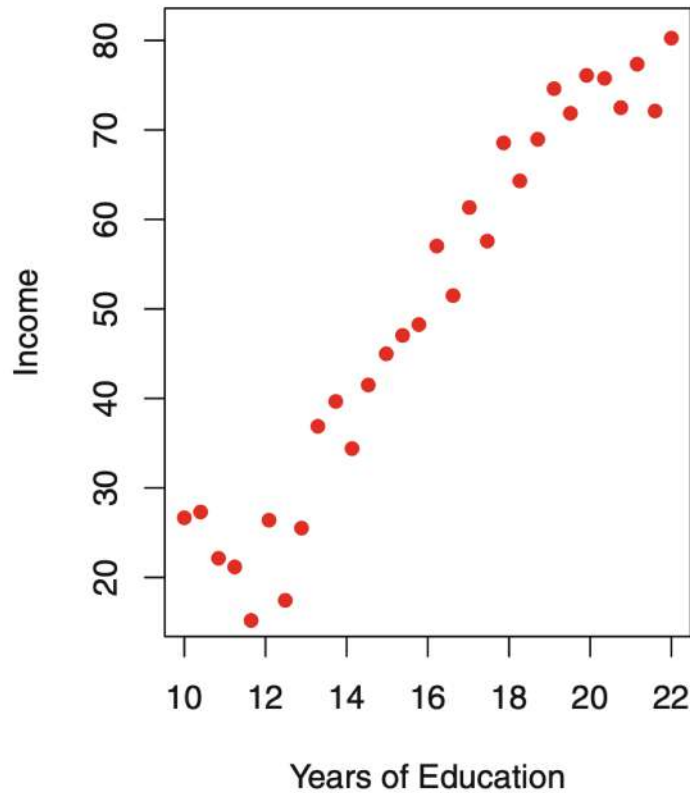
Observations
(simulated)



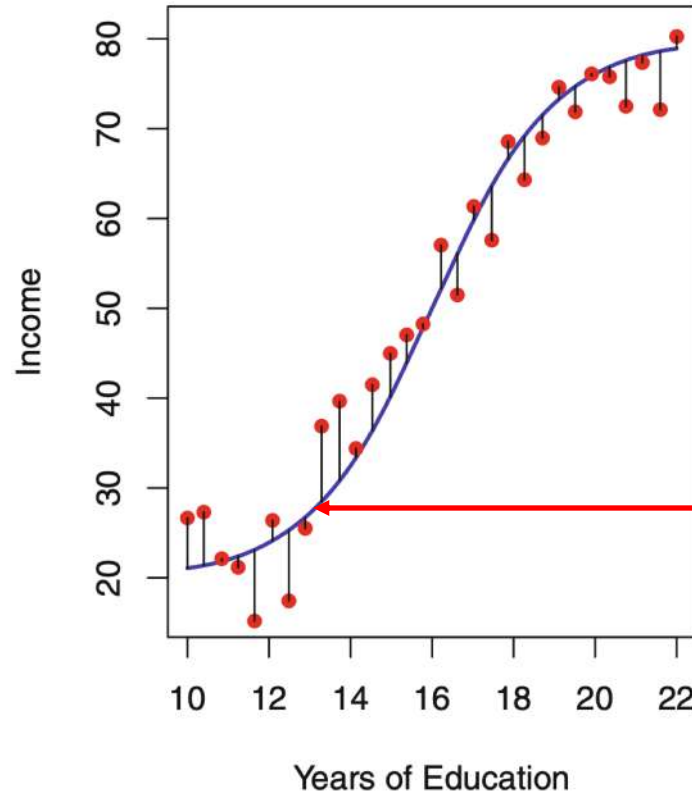
True yet unknown relationship
between X and Y

$$Y = f(X) + \varepsilon$$

Example: Y =Income vs. X =Education



Observations
(simulated)

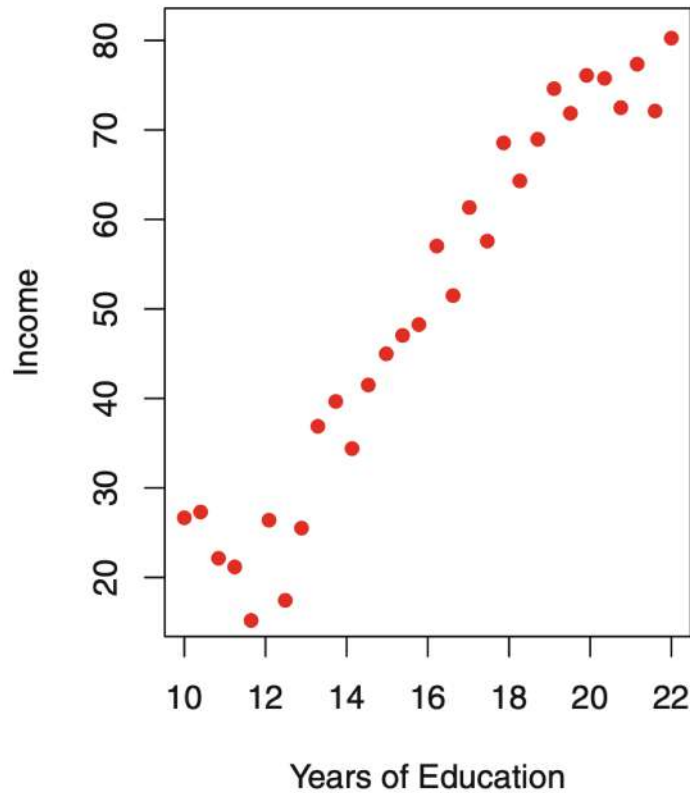


True yet unknown relationship
between X and Y

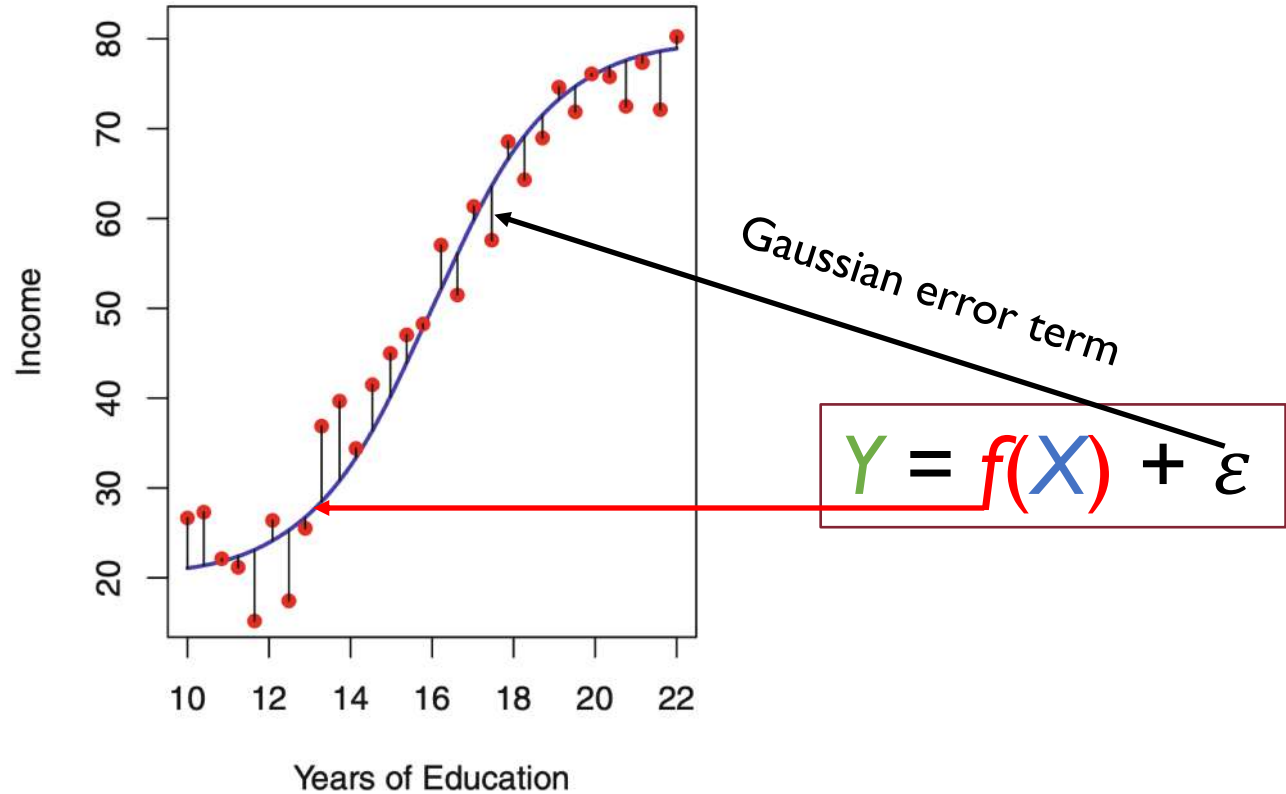
$$Y = f(X) + \varepsilon$$

A red arrow points from the $f(X)$ term in the equation to the blue curve in the adjacent plot.

Example: Y =Income vs. X =Education



Observations
(simulated)



True yet unknown relationship
between X and Y

Assumptions

- There exists a relationship between \mathcal{X} (features) and \mathcal{Y} (values)

$$\mathcal{Y} = f(\mathcal{X}) + \epsilon$$

Assumptions

- There exists a relationship between X (features) and Y (values)

$$\mathcal{Y} = f(\mathcal{X}) + \epsilon$$

- f is some fixed but unknown function of X

Assumptions

- There exists a relationship between X (features) and Y (values)

$$\mathcal{Y} = f(\mathcal{X}) + \epsilon$$

- f is some fixed but unknown function of X
- ϵ is a random error term, which is independent of X and has 0-mean

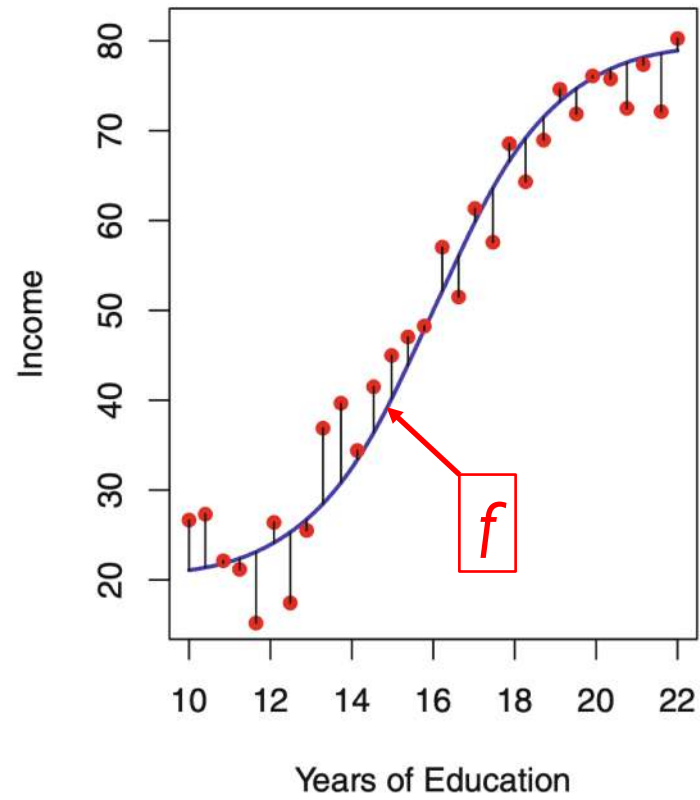
Assumptions

- There exists a relationship between X (features) and Y (values)

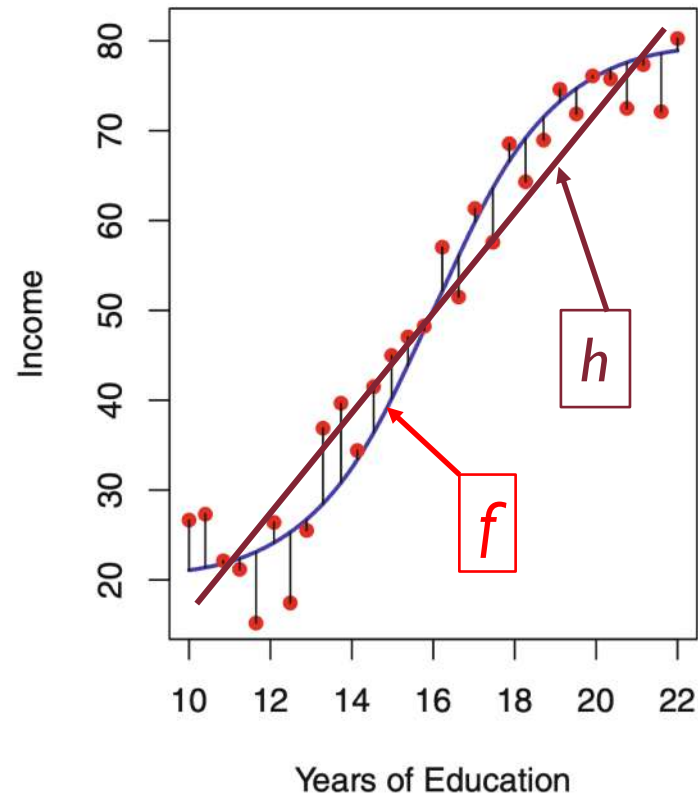
$$\mathcal{Y} = f(\mathcal{X}) + \epsilon$$

- f is some fixed but unknown function of X
- ϵ is a random error term, which is independent of X and has 0-mean
- In this formulation, f represents the systematic information that X provides about Y

Goal

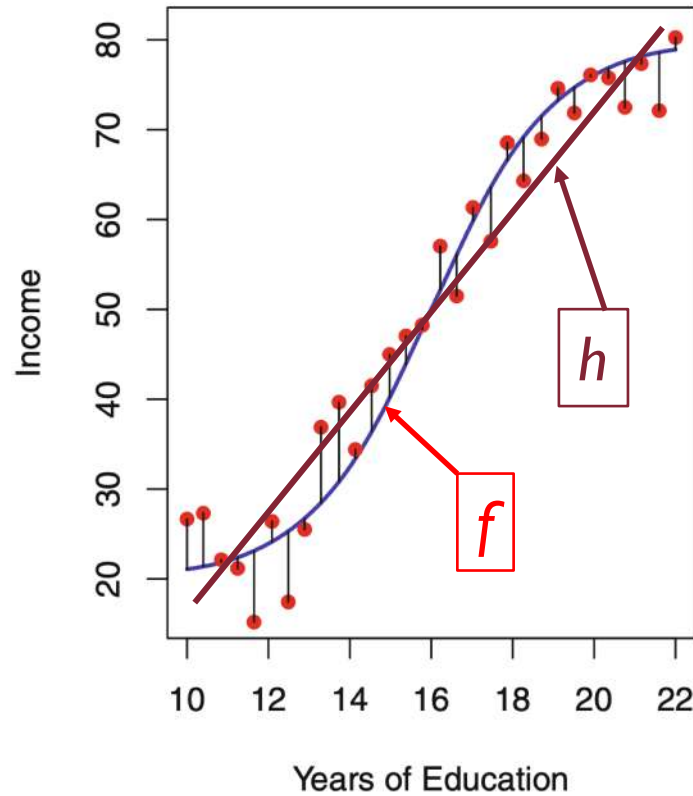


Goal



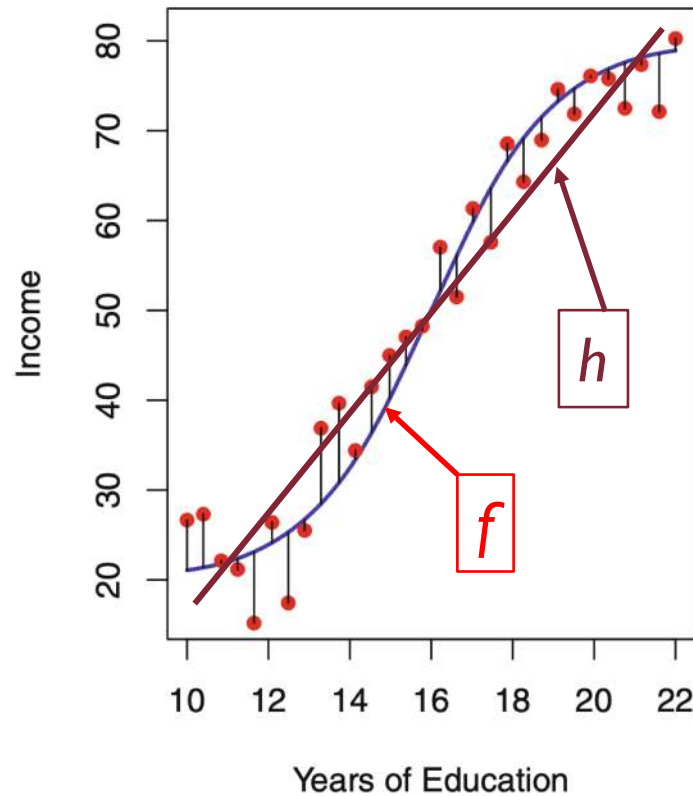
- Find an approximation h of the true relationship f

Goal



- Find an approximation h of the true relationship f
- Choose h from a specific hypothesis space H (i.e., linear functions)

Goal



- Find an approximation h of the true relationship f
- Choose h from a specific hypothesis space H (i.e., linear functions)
- Use a dataset D of observations to learn h

$$h(X) \sim f(X)$$

Recap of Notation

$$\mathcal{X} \subseteq \mathbb{R}^n$$

input feature space

$$\mathcal{Y}$$

output space

$$y \in \mathbb{R}$$

real-value label of the i -th instance
(**regression**)

$$(\mathbf{x}_i, y_i)$$

i -th labeled instance

$$\mathbf{x}_i = (x_{i,1}, \dots, x_{i,n}) \in \mathcal{X}$$

n -dimensional feature vector of the i -th instance

$$y_i \in \mathcal{Y}$$

label of the i -th instance

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)\}$$

dataset of m i.i.d. labeled instances

The Hypothesis Space H

The hypothesis space is defined as follows:

$$\mathcal{H} = \{h_{\theta} : \mathcal{X} \mapsto \mathcal{Y} \mid h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n\}$$

The Hypothesis Space H

The hypothesis space is defined as follows:

$$\mathcal{H} = \{h_{\theta} : \mathcal{X} \mapsto \mathcal{Y} \mid h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n\}$$

θ $n+1$ -dimensional vector of model parameters

The Hypothesis Space H

The hypothesis space is defined as follows:

$$\mathcal{H} = \{h_{\theta} : \mathcal{X} \mapsto \mathcal{Y} \mid h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n\}$$

θ $n+1$ -dimensional vector of model parameters

$x_0 = 1$ by convention

The Hypothesis Space H

The hypothesis space is defined as follows:

$$\mathcal{H} = \{h_{\theta} : \mathcal{X} \mapsto \mathcal{Y} \mid h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n\}$$

θ $n+1$ -dimensional vector of model parameters

$x_0 = 1$ by convention

Among all the possible instantiations of θ the learning algorithm selects θ^* as the one which minimizes a **loss function** measured on D

Residuals of Sum Squares (RSS)

$$y_i = f(\mathbf{x}_i) + \epsilon_i \quad i\text{-th observation}$$

Residuals of Sum Squares (RSS)

$$y_i = f(\mathbf{x}_i) + \epsilon_i \quad i\text{-th observation} \qquad \hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i) \quad i\text{-th prediction}$$

Residuals of Sum Squares (RSS)

$$y_i = f(\mathbf{x}_i) + \epsilon_i \quad i\text{-th observation} \qquad \hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i) \quad i\text{-th prediction}$$

$$\hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) = \theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}$$

Residuals of Sum Squares (RSS)

$$y_i = f(\mathbf{x}_i) + \epsilon_i \quad i\text{-th observation} \qquad \hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i) \quad i\text{-th prediction}$$

$$\hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) = \theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}$$

$$e_i = \hat{y}_i - y_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) - \underbrace{y_i}_{f(\mathbf{x}_i) + \epsilon_i} \quad i\text{-th residual}$$

Residuals of Sum Squares (RSS)

$$y_i = f(\mathbf{x}_i) + \epsilon_i \quad i\text{-th observation} \qquad \hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i) \quad i\text{-th prediction}$$

$$\hat{y}_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) = \theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}$$

$$e_i = \hat{y}_i - y_i = h_{\boldsymbol{\theta}}(\mathbf{x}_i) - \underbrace{y_i}_{f(\mathbf{x}_i) + \epsilon_i} \quad i\text{-th residual}$$

$$\text{RSS}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \sum_{i=1}^m e_i^2 = \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2$$

Ordinary Least Squares (OLS)

- Remember that the supervised learning problem can be generally defined as the following optimization problem

$$h^* = \operatorname{argmin}_{h \in \mathcal{H}} L(h, \mathcal{D})$$

Ordinary Least Squares (OLS)

- Remember that the supervised learning problem can be generally defined as the following optimization problem

$$h^* = \operatorname{argmin}_{h \in \mathcal{H}} L(h, \mathcal{D})$$

- OLS is the usual approach to fit (i.e., find the optimal set of parameters of) linear regression models

Ordinary Least Squares (OLS)

- Remember that the supervised learning problem can be generally defined as the following optimization problem

$$h^* = \operatorname{argmin}_{h \in \mathcal{H}} L(h, \mathcal{D})$$

- OLS is the usual approach to fit (i.e., find the optimal set of parameters of) linear regression models

$$h^* = h_{\theta^*} = \operatorname{argmin}_{\theta} L(h_{\theta}, \mathcal{D})$$

The Loss Function L : Mean Squared Error

- OLS uses **Mean Squared Error** as the loss function to minimize

The Loss Function L : Mean Squared Error

- OLS uses **Mean Squared Error** as the loss function to minimize
- MSE measures the **average error** when the true f is substituted with a hypothesis h_θ in H (in-sample error)

The Loss Function L : Mean Squared Error

- OLS uses **Mean Squared Error** as the loss function to minimize
- MSE measures the **average error** when the true f is substituted with a hypothesis h_{θ} in H (in-sample error)

$$\begin{aligned}\text{MSE}(h_{\theta}, \mathcal{D}) &= \frac{1}{m} \text{RSS}(h_{\theta}, \mathcal{D}) = \\ &= \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2\end{aligned}$$

The OLS Learning Algorithm

OLS aims at solving the following optimization problem:

$$\begin{aligned} h^* = h_{\theta^*} &= \operatorname{argmin}_{\theta} \operatorname{MSE}(h_{\theta}, \mathcal{D}) = \\ &= \operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right] \end{aligned}$$

The OLS Learning Algorithm

OLS aims at solving the following optimization problem:

$$\begin{aligned} h^* &= h_{\theta^*} = \operatorname{argmin}_{\theta} \operatorname{MSE}(h_{\theta}, \mathcal{D}) = \\ &= \operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right] \end{aligned}$$

How do we solve that?

Min/Max of a Convex/Concave Function

- Any local minimum (maximum) of a convex (concave) function is also a global minimum (maximum)

Min/Max of a Convex/Concave Function

- Any local minimum (maximum) of a convex (concave) function is also a global minimum (maximum)
- If the function is convex (concave) finding the minimum (maximum) can be done just by computing the first derivative and set it to 0

Min/Max of a Convex/Concave Function

- Any local minimum (maximum) of a convex (concave) function is also a global minimum (maximum)
- If the function is convex (concave) finding the minimum (maximum) can be done just by computing the first derivative and set it to 0
- In the case of a multivariate function, this generalizes to compute the gradient (∇) of the function and set it to 0

The Gradient ∇

The gradient of an n -variable function is the n -dimensional vector of the **partial derivatives** of the function w.r.t. each of its variable

The Gradient ∇

The gradient of an n -variable function is the n -dimensional vector of the **partial derivatives** of the function w.r.t. each of its variable

$$f : \mathbb{R}^n \mapsto \mathbb{R} \quad \nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$

The Gradient ∇

The gradient of an n -variable function is the n -dimensional vector of the **partial derivatives** of the function w.r.t. each of its variable

$$f : \mathbb{R}^n \mapsto \mathbb{R} \quad \nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$

Solving $\nabla f = \mathbf{0}$ means finding the n -dimensional vector \mathbf{x} such that:

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right) = \underbrace{(0, 0, \dots, 0)}_n = \mathbf{0}$$

Solving the Optimization Problem

$$\operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

Solving the Optimization Problem

$$\operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right]$$

Observations y_i and features \mathbf{x}_i can be thought of as fixed constants

Solving the Optimization Problem

$$\operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right]$$

Observations y_i and features \mathbf{x}_i can be thought of as fixed constants

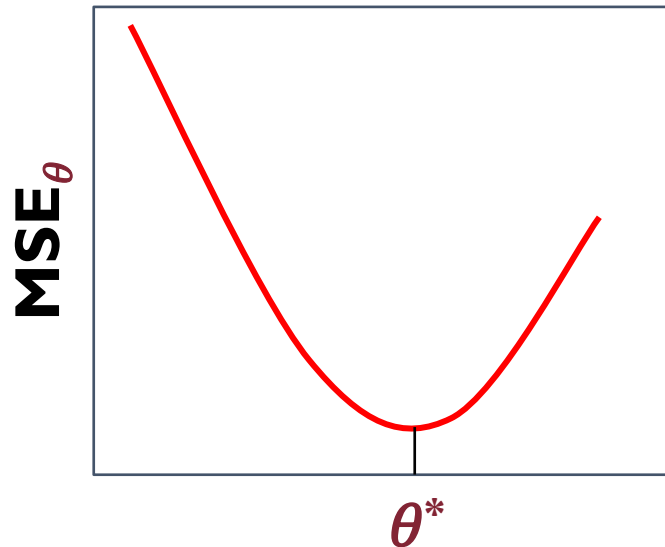
Each term of the summation is a multivariate linear function of the model parameters θ

Solving the Optimization Problem

$$\operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right]$$

Observations y_i and features \mathbf{x}_i can be thought of as fixed constants

Each term of the summation is a multivariate linear function of the model parameters θ



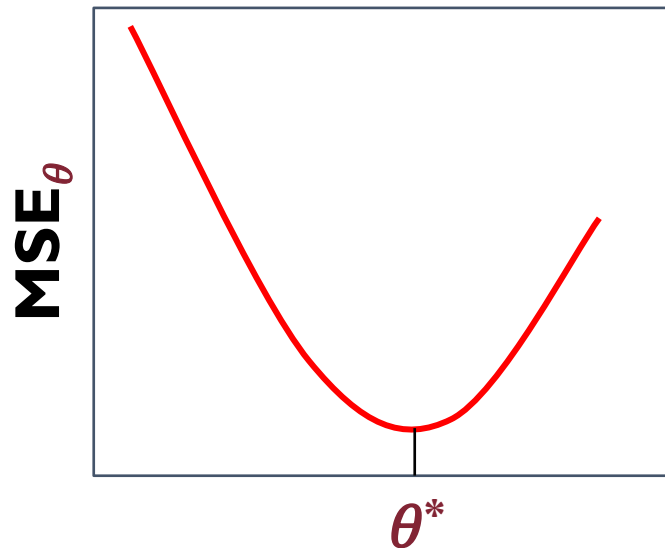
Linear functions are convex and so do sum of those

Solving the Optimization Problem

$$\operatorname{argmin}_{\theta} \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i)^2 \right]$$

Observations y_i and features \mathbf{x}_i can be thought of as fixed constants

Each term of the summation is a multivariate linear function of the model parameters θ



Linear functions are convex and so do sum of those

Convex functions have a **unique local minimum**, which therefore happens to be the **global minimum**

Computing the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

Computing the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

$$\frac{\partial f}{\partial t}(\alpha t) = \alpha \frac{\partial f}{\partial t}(t) \quad \alpha \in \mathbb{R}, \text{ constant}$$

scalar multiple rule

Computing the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

$$\frac{\partial f}{\partial t}(\alpha t) = \alpha \frac{\partial f}{\partial t}(t) \quad \alpha \in \mathbb{R}, \text{ constant}$$

scalar multiple rule

$$\frac{\partial f}{\partial t} \left(\sum t \right) = \left(\sum \frac{\partial f}{\partial t}(t) \right)$$

sum rule

Computing the Gradient of MSE

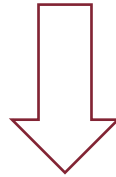
$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

$$\frac{\partial f}{\partial t}(\alpha t) = \alpha \frac{\partial f}{\partial t}(t) \quad \alpha \in \mathbb{R}, \text{ constant}$$

scalar multiple rule

$$\frac{\partial f}{\partial t} \left(\sum t \right) = \left(\sum \frac{\partial f}{\partial t}(t) \right)$$

sum rule



$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{1}{m} \left[\sum_{i=1}^m \nabla (h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)^2 \right]$$

Computing the Gradient of MSE (1 instance)

To make things easier, let's assume the dataset D contains a single instance (\mathbf{x}, y)

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla (h_{\theta}(\mathbf{x}) - y)^2$$

Computing the Gradient of MSE (1 instance)

To make things easier, let's assume the dataset D contains a single instance (\mathbf{x}, y)

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla (h_{\theta}(\mathbf{x}) - y)^2$$

$$\frac{\partial f}{\partial t} (t)^{\alpha} = \alpha t^{\alpha-1}$$

power rule

Computing the Gradient of MSE (1 instance)

To make things easier, let's assume the dataset D contains a single instance (\mathbf{x}, y)

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla (h_{\theta}(\mathbf{x}) - y)^2$$

$$\frac{\partial f}{\partial t}(t)^{\alpha} = \alpha t^{\alpha-1}$$

power rule

$$\frac{\partial f}{\partial t}(g(t)) = \frac{\partial f}{\partial g(t)}(g(t)) \frac{\partial g}{\partial t}(t)$$

chain rule

Computing the Gradient of MSE (1 instance)

To make things easier, let's assume the dataset D contains a single instance (\mathbf{x}, y)

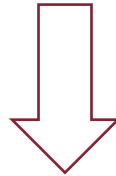
$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla (h_{\theta}(\mathbf{x}) - y)^2$$

$$\frac{\partial f}{\partial t}(t)^{\alpha} = \alpha t^{\alpha-1}$$

power rule

$$\frac{\partial f}{\partial t}(g(t)) = \frac{\partial f}{\partial g(t)}(g(t)) \frac{\partial g}{\partial t}(t)$$

chain rule



$$2(h_{\theta}(\mathbf{x}) - y) \nabla (h_{\theta}(\mathbf{x}) - y)$$

Computing the Gradient of MSE (1 instance)

To make things easier, let's assume the dataset D contains a single instance (\mathbf{x}, y)

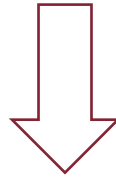
$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla (h_{\theta}(\mathbf{x}) - y)^2$$

$$\frac{\partial f}{\partial t}(t)^{\alpha} = \alpha t^{\alpha-1}$$

power rule

$$\frac{\partial f}{\partial t}(g(t)) = \frac{\partial f}{\partial g(t)}(g(t)) \frac{\partial g}{\partial t}(t)$$

chain rule



$$2(h_{\theta}(\mathbf{x}) - y) \nabla(h_{\theta}(\mathbf{x}) - y)$$

Computing the Gradient of MSE (1 instance)

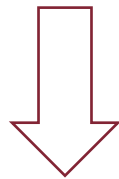
$$\nabla(h_{\boldsymbol{\theta}}(\mathbf{x}) - y) = \nabla(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y) =$$

Computing the Gradient of MSE (1 instance)

$$\begin{aligned}\nabla(h_{\boldsymbol{\theta}}(\mathbf{x}) - y) &= \nabla(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y) = \\ &= \underbrace{\left(\frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_0}, \dots, \frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_n} \right)}_{n+1} = (x_0, x_1, \dots, x_n) = \mathbf{x}\end{aligned}$$

Computing the Gradient of MSE (1 instance)

$$\begin{aligned}\nabla(h_{\boldsymbol{\theta}}(\mathbf{x}) - y) &= \nabla(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y) = \\ &= \underbrace{\left(\frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_0}, \dots, \frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_n} \right)}_{n+1} = (x_0, x_1, \dots, x_n) = \mathbf{x}\end{aligned}$$



$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \underbrace{2(h_{\boldsymbol{\theta}}(\mathbf{x}) - y)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}}_{(n+1)\text{-dimensional vector}}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \underbrace{2(h_{\boldsymbol{\theta}}(\mathbf{x}) - y)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}}_{(n+1)\text{-dimensional vector}}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \underbrace{2(h_{\boldsymbol{\theta}}(\mathbf{x}) - y)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}}_{(n+1)\text{-dimensional vector}}$$

$$h_{\boldsymbol{\theta}}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n = \boldsymbol{\theta}^T \cdot \mathbf{x}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \underbrace{2(h_{\boldsymbol{\theta}}(\mathbf{x}) - y)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}}_{(n+1)\text{-dimensional vector}}$$

$$h_{\boldsymbol{\theta}}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n = \boldsymbol{\theta}^T \cdot \mathbf{x}$$



$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)\mathbf{x}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)\mathbf{x}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)\mathbf{x}$$

$$= \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_0 \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix} = \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix}$$

Vectorized Notation

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)\mathbf{x}$$

$$= \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_0 \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix} = \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix}$$

$x_0 = 1$ by definition

Vectorized Notation

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)\mathbf{x}$$

$$= \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_0 \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix} = \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix}$$

$x_0 = 1$ by definition

The resulting gradient is an $(n+1)$ -dimensional vector as expected!

Setting the Gradient Equal to Zero

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}$$

Setting the Gradient Equal to Zero

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \begin{bmatrix} 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_n \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}$$

We need to solve a system of $n+1$ linear equations with $n+1$ variables

$$2(\boldsymbol{\theta}^T \cdot \mathbf{x} - y)x_j = 0 \quad \forall j \in \{0, 1, \dots, n\}$$

Computing the Gradient of MSE (m instances)

In the general case where the dataset \mathcal{D} contains a m instances

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \nabla \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \right]$$

Computing the Gradient of MSE (m instances)

In the general case where the dataset \mathcal{D} contains a m instances

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \nabla \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \right]$$



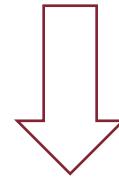
$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \underbrace{\left(h_{\theta}(\mathbf{x}_i) - y_i \right)}_{\text{scalar}} \underbrace{\mathbf{x}_i}_{n+1\text{-dimensional vector}} \right]$$

Computing the Gradient of MSE (m instances)

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \underbrace{\left(h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)}_{\text{scalar}} \underbrace{\mathbf{x}_i}_{n+1\text{-dimensional vector}} \right]$$

Computing the Gradient of MSE (m instances)

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \underbrace{\left(h_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i \right)}_{\text{scalar}} \underbrace{\mathbf{x}_i}_{n+1\text{-dimensional vector}} \right]$$



$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \begin{bmatrix} \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,0} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,0} \\ \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,1} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,1} \\ \vdots \\ \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,n} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,n} \end{bmatrix}$$

Computing the Gradient of MSE (m instances)

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \begin{bmatrix} (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1) + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m) \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,1} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,1} \\ \vdots \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,n} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,n} \end{bmatrix}$$

Computing the Gradient of MSE (m instances)

$$x_{i,0} = 1 \text{ for all } i = \{1, \dots, m\}$$

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \begin{bmatrix} (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1) + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m) \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,1} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,1} \\ \vdots \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,n} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,n} \end{bmatrix}$$

Computing the Gradient of MSE (m instances)

$$x_{i,0} = 1 \text{ for all } i = \{1, \dots, m\}$$

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \begin{bmatrix} (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1) + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m) \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,1} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,1} \\ \vdots \\ (\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,n} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,n} \end{bmatrix}$$

Again, we need to solve a system of $n+1$ linear equations with $n+1$ variables

$$\frac{2}{m} \left[(\boldsymbol{\theta}^T \cdot \mathbf{x}_1 - y_1)x_{1,j} + \dots + (\boldsymbol{\theta}^T \cdot \mathbf{x}_m - y_m)x_{m,j} \right] = 0 \quad \forall j \in \{0, \dots, n\}$$

Matrix Notation

$$\mathbf{X} = \underbrace{\begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n} \\ \vdots & \vdots & \vdots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n} \end{bmatrix}}_{m \times n+1 \text{ feature matrix}} = \begin{bmatrix} -\mathbf{x}_1^T - \\ -\mathbf{x}_2^T - \\ \vdots \\ -\mathbf{x}_m^T - \end{bmatrix}$$

Matrix Notation

$$\mathbf{X} = \underbrace{\begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n} \\ \vdots & \vdots & \vdots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n} \end{bmatrix}}_{m \times n+1 \text{ feature matrix}} = \begin{bmatrix} -\mathbf{x}_1^T - \\ -\mathbf{x}_2^T - \\ \vdots \\ -\mathbf{x}_m^T - \end{bmatrix}$$

$$\boldsymbol{\theta} = \underbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}}$$

$n+1$ -dimensional parameter vector

Matrix Notation

$$\mathbf{X} = \underbrace{\begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n} \\ \vdots & \vdots & \vdots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n} \end{bmatrix}}_{m \times n+1 \text{ feature matrix}} = \begin{bmatrix} -\mathbf{x}_1^T - \\ -\mathbf{x}_2^T - \\ \vdots \\ -\mathbf{x}_m^T - \end{bmatrix}$$

$$\boldsymbol{\theta} = \underbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}}$$

$n+1$ -dimensional **parameter vector**

$$\mathbf{y} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}}$$

m -dimensional **target vector**

Vectorized Form of the Optimization Problem

$$h^* = h_{\theta^*} = \operatorname{argmin}_{\theta} \left[\underbrace{\frac{1}{m} \|\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}\|^2}_{\text{MSE}(h_{\theta}, \mathcal{D})} \right]$$

Vectorized Form of the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y})$$

Vectorized Form of the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y})$$

$$\frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

Vectorized Form of the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y})$$

$$\frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \mathbf{X} \cdot \boldsymbol{\theta} = \mathbf{X}^T \cdot \mathbf{y}$$

Vectorized Form of the Gradient of MSE

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = \frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y})$$

$$\frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \mathbf{X} \cdot \boldsymbol{\theta} = \mathbf{X}^T \cdot \mathbf{y}$$

$$\boldsymbol{\theta} = \mathbf{X}^\dagger \cdot \mathbf{y}$$

$\mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$ is the **pseudo-inverse** of \mathbf{X}

The Pseudo-Inverse of \mathbf{X}

- In general, the feature matrix \mathbf{X} is non-squared therefore non-invertible

The Pseudo-Inverse of \mathbf{X}

- In general, the feature matrix \mathbf{X} is non-squared therefore non-invertible
- $\mathbf{X}^T\mathbf{X}$ is instead square (n -by- n) and **very likely** invertible
 - The chance of a randomly generated squared matrix is invertible approaches 1
 - To be non-invertible, the determinant must be 0 (linearly dependent columns)

The Pseudo-Inverse of \mathbf{X}

- In general, the feature matrix \mathbf{X} is non-squared therefore non-invertible
- $\mathbf{X}^T\mathbf{X}$ is instead square (n -by- n) and **very likely** invertible
 - The chance of a randomly generated squared matrix is invertible approaches 1
 - To be non-invertible, the determinant must be 0 (linearly dependent columns)
- Typically, the number m of rows (instances) are way larger than the number n of columns (features)
 - $\mathbf{X}^T\mathbf{X}$ is smaller than \mathbf{X}

Additional Notes on OLS

- OLS is also known as one-step learning as there exists a closed-form (i.e., analytical) solution to the convex optimization problem

Additional Notes on OLS

- OLS is also known as one-step learning as there exists a closed-form (i.e., analytical) solution to the convex optimization problem
- However, other choices of loss functions (even if convex) may need an **iterative** approach to get to a (local) minimum

Additional Notes on OLS

- OLS is also known as one-step learning as there exists a closed-form (i.e., analytical) solution to the convex optimization problem
- However, other choices of loss functions (even if convex) may need an **iterative** approach to get to a (local) minimum
- Though in general $n \ll m$, computing the inverse of an n -by- n matrix is still a costly operation ($O(n^3)$ time complexity)

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th observation

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th observation

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

i -th unobservable **error**

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th observation

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

i -th unobservable **error**

i -th prediction

$$\hat{y}_i = h_{\theta}(\mathbf{x}_i) \approx f(\mathbf{x}_i)$$

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th observation

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

i -th unobservable **error**

i -th prediction

$$\hat{y}_i = h_{\theta}(\mathbf{x}_i) \approx f(\mathbf{x}_i)$$

$$e_i = \hat{y}_i - y_i = h_{\theta}(\mathbf{x}_i) - \underbrace{y_i}_{f(\mathbf{x}_i) + \epsilon_i}$$

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th observation

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

i -th unobservable **error**

i -th prediction

$$\hat{y}_i = h_{\theta}(\mathbf{x}_i) \approx f(\mathbf{x}_i)$$

$$e_i = \hat{y}_i - y_i = h_{\theta}(\mathbf{x}_i) - \underbrace{y_i}_{f(\mathbf{x}_i) + \epsilon_i}$$

i -th residual

Errors vs. Residuals

Subtle yet important difference between **errors** and **residuals**

i -th **observation**

$$y_i = f(\mathbf{x}_i) + \epsilon_i$$

i -th unobservable **error**

i -th **prediction**

$$\hat{y}_i = h_{\theta}(\mathbf{x}_i) \approx f(\mathbf{x}_i)$$

$$\underbrace{e_i}_{i\text{-th residual}} = \hat{y}_i - y_i = h_{\theta}(\mathbf{x}_i) - \underbrace{y_i}_{f(\mathbf{x}_i) + \epsilon_i}$$

MSE is computed from residuals, not unobservable errors!

OLS Assumptions

- **Weak exogeneity** → Predictor variables (i.e., features) can be treated as error-free constants

OLS Assumptions

- **Weak exogeneity** → Predictor variables (i.e., features) can be treated as error-free constants
- **Linearity** → Linear relationship between the features and the response
 - Only a restriction on the parameters; features themselves can be arbitrarily combined using non-linear transformations

OLS Assumptions

- **Weak exogeneity** → Predictor variables (i.e., features) can be treated as error-free constants
- **Linearity** → Linear relationship between the features and the response
 - Only a restriction on the parameters; features themselves can be arbitrarily combined using non-linear transformations
- **Error independence** → Error terms ε_i are uncorrelated with each other
 - Knowing that ε_i is positive (negative) gives no information on the sign of ε_{i+1}

OLS Assumptions

- **Homoscedasticity** → Different values of the response variable have the same variance in their errors, regardless of the feature values
 - In practice, this does not hold when the response varies over a wide scale

OLS Assumptions

- **Homoscedasticity** → Different values of the response variable have the same variance in their errors, regardless of the feature values
 - In practice, this does not hold when the response varies over a wide scale
- **No Multicollinearity** → There must not be two or more features whose values are perfectly correlated with each other
 - The feature matrix \mathbf{X} must have full column rank n
 - If \mathbf{X} is full column rank n then $\mathbf{X}^T\mathbf{X}$ is always invertible
 - It can be shown that if $\mathbf{X}^T\mathbf{X}\mathbf{u} = \mathbf{0}$ for some vector \mathbf{u} , then $\mathbf{u} = \mathbf{0}$ (trivial solution)

Checking OLS Assumptions

- A good way to assess the OLS assumptions hold is to use residual plots

Checking OLS Assumptions

- A good way to assess the OLS assumptions hold is to use residual plots
- Plotting residuals against each feature and/or the predicted value may help spot:
 - Non-linearity
 - Correlation between error terms
 - Non-constant variance of error terms (i.e., heteroscedasticity)
 - ...

Assessing the Quality of a Model

- Suppose we have fit a linear regression model to some dataset of observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1..m}$

Assessing the Quality of a Model

- Suppose we have fit a linear regression model to some dataset of observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1..m}$
- In other words, we estimated the vector of parameters θ^* using OLS

Assessing the Quality of a Model

- Suppose we have fit a linear regression model to some dataset of observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1..m}$
- In other words, we estimated the vector of parameters θ^* using OLS

How do we measure the "goodness-of-fit" of the model?

Assessing the Quality of a Model

- Suppose we have fit a linear regression model to some dataset of observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1..m}$
- In other words, we estimated the vector of parameters θ^* using OLS

How do we measure the "goodness-of-fit" of the model?



Residual Standard Error
(RSE)

Assessing the Quality of a Model

- Suppose we have fit a linear regression model to some dataset of observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1..m}$
- In other words, we estimated the vector of parameters θ^* using OLS

How do we measure the "goodness-of-fit" of the model?

Residual Standard Error
(RSE)

R^2 statistic

Residual Standard Error (RSE)

Recall that every observation of the target variable y_i is associated with an error term ϵ_i

$$y_i = \underbrace{\theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}}_{h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i)} + \epsilon_i$$

Residual Standard Error (RSE)

Recall that every observation of the target variable y_i is associated with an error term ϵ_i

$$y_i = \underbrace{\theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}}_{h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i)} + \boxed{\epsilon_i}$$

Residual Standard Error (RSE)

Recall that every observation of the target variable y_i is associated with an error term ε_i

$$y_i = \underbrace{\theta_0 x_{i,0} + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n}}_{h_{\boldsymbol{\theta}}(\mathbf{x}_i) \approx f(\mathbf{x}_i)} + \boxed{\epsilon_i}$$

Even if we were able to find the exact parameters of the true f , we would not be able to perfectly predict y_i from \mathbf{x}_i

Residual Standard Error (RSE)

RSE is an estimate of the standard deviation of ε

$$\text{RSE}(h_{\theta}, \mathcal{D}) = \sqrt{\frac{1}{\underbrace{m - n - 1}_{\text{degrees of freedom}}} \underbrace{\sum_{i=1}^m (\hat{y}_i - y_i)^2}_{\text{RSS}}}$$

Residual Standard Error (RSE)

RSE is an estimate of the standard deviation of ε

$$\text{RSE}(h_{\theta}, \mathcal{D}) = \sqrt{\frac{1}{\underbrace{m - n - 1}_{\text{degrees of freedom}}} \underbrace{\sum_{i=1}^m (\hat{y}_i - y_i)^2}_{\text{RSS}}}$$

A measure of the **lack** of fit of the model to the data
the lower the better

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$

Degrees of Freedom


$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$

How many observations n do I need to estimate model's parameters?

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$


How many observations m do I need to estimate model's parameters?



Does $m = 1$ sound enough to fit a line?

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$




How many observations n do I need to estimate model's parameters?

With 2 data points I am always able to fit a perfect line

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$




How many observations m do I need to estimate model's parameters?

With 2 data points I am always able to fit a perfect line

Problem is that my fitted line may drastically change depending on where the second point is located!

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$




How many observations n do I need to estimate model's parameters?

With 2 data points I am always able to fit a perfect line

If I want my model to be more "flexible" I need **at least** 3 points which leave me with 1 degree of freedom

Degrees of Freedom

$$y_i = \theta_0 + \theta_1 x_{i,1} + \epsilon_i$$



How many observations n do I need to estimate model's parameters?

With 2 data points I am always able to fit a perfect line

If I want my model to be more "flexible" I need **at least** 3 points which leave me with 1 degree of freedom

Degrees of Freedom

What happens when we add more variables to the model?

$$y_i = \theta_0 + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n} + \epsilon_i$$

Degrees of Freedom

What happens when we add more variables to the model?

$$y_i = \theta_0 + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n} + \epsilon_i$$

Given the same number of observations m ,
we lose 1 degree of freedom for each variable we add

Degrees of Freedom

What happens when we add more variables to the model?

$$y_i = \theta_0 + \theta_1 x_{i,1} + \dots + \theta_n x_{i,n} + \epsilon_i$$

Given the same number of observations m ,
we lose 1 degree of freedom for each variable we add

$$\text{df} = \underbrace{m}_{\text{\#observations}} - \underbrace{n}_{\text{\#features}} - \underbrace{1}_{\text{intercept}}$$

R² Statistic

$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}} = 1 - \frac{\sum_{i=1}^m (\hat{y}_i - y_i)^2}{\sum_{i=1}^m (y_i - \bar{y})^2}$$

TSS measures the total variance in the response **Y** before the regression takes place

R² Statistic

$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}} = 1 - \frac{\sum_{i=1}^m (\hat{y}_i - y_i)^2}{\sum_{i=1}^m (y_i - \bar{y})^2}$$

TSS measures the total variance in the response **Y** before the regression takes place

RSS measures the amount of variability that is left unexplained after performing the regression

R² Statistic

$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}} = 1 - \frac{\sum_{i=1}^m (\hat{y}_i - y_i)^2}{\sum_{i=1}^m (y_i - \bar{y})^2}$$

TSS measures the total variance in the response **Y** before the regression takes place

RSS measures the amount of variability that is left unexplained after performing the regression

R² measures the proportion of variability in **Y** that can be explained using **X**

R^2 Statistic

- An R^2 statistic that is close to 1 indicates that a large proportion of the variability in the response has been explained by the regression

R^2 Statistic

- An R^2 statistic that is close to 1 indicates that a large proportion of the variability in the response has been explained by the regression
- The larger R^2 the better is the linear regression model

R^2 Statistic

- An R^2 statistic that is close to 1 indicates that a large proportion of the variability in the response has been explained by the regression
- The larger R^2 the better is the linear regression model
- R^2 is easier to interpret than RSE as it always ranges between 0 and 1

Adjusted R^2 Statistic

- Fixing the sample size m , RSS decreases (or, at worst, it stays the same) as more variables are added to the fitted model

Adjusted R^2 Statistic

- Fixing the sample size m , RSS decreases (or, at worst, it stays the same) as more variables are added to the fitted model
- R^2 always increases as more variables are added (as df decreases!)

Adjusted R^2 Statistic

- Fixing the sample size m , RSS decreases (or, at worst, it stays the same) as more variables are added to the fitted model
- R^2 always increases as more variables are added (as df decreases!)
- We need a way to adjust for that, otherwise we could get a better model by simply adding useless features to it!

$$R_{\text{adj}}^2 = 1 - \frac{\frac{\text{RSS}}{m-n-1}}{\frac{\text{TSS}}{m-1}}$$

Adjusted R^2 Statistic

- Maximizing the adjusted R^2 is equivalent to minimizing $RSS/(m-n-1)$

Adjusted R^2 Statistic

- Maximizing the adjusted R^2 is equivalent to minimizing $RSS/(m-n-1)$
- We know RSS may decrease if the number of variables in the model increases

Adjusted R^2 Statistic

- Maximizing the adjusted R^2 is equivalent to minimizing $RSS/(m-n-1)$
- We know RSS may decrease if the number of variables in the model increases
- $RSS/(m-n-1)$ may increase or decrease, due to the presence of n in the denominator

Adjusted R^2 Statistic

- Maximizing the adjusted R^2 is equivalent to minimizing $RSS/(m-n-1)$
- We know RSS may decrease if the number of variables in the model increases
- $RSS/(m-n-1)$ may increase or decrease, due to the presence of n in the denominator
- We may need to increase the sample size m to compensate for the increasing of RSS due to the inclusion of more features n

Regularization

- The absolute value of learned parameters θ should not be very large
 - Otherwise, a small change in an input feature may cause a high difference in the output predicted value

Regularization

- The absolute value of learned parameters θ should not be very large
 - Otherwise, a small change in an input feature may cause a high difference in the output predicted value
- This is an indication of overfitting:
 - The learned model is highly "training set dependent" and does not generalize

Regularization

- The absolute value of learned parameters θ should not be very large
 - Otherwise, a small change in an input feature may cause a high difference in the output predicted value
- This is an indication of overfitting:
 - The learned model is highly "training set dependent" and does not generalize
- **Regularization** → Put some constraint on the optimization problem so as to limit the values of the learned parameters

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} ||\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}||^2 + \lambda \left(\alpha ||\boldsymbol{\theta}|| + (1 - \alpha) ||\boldsymbol{\theta}||^2 \right) \right]$$

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} ||\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}||^2 + \lambda \left(\alpha ||\boldsymbol{\theta}|| + (1 - \alpha) ||\boldsymbol{\theta}||^2 \right) \right]$$

$\lambda \geq 0$ **regularization parameter:** when this is 0 we backup to OLS (no regularization at all)

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} \|\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}\|^2 + \lambda \left(\alpha \|\boldsymbol{\theta}\| + (1 - \alpha) \|\boldsymbol{\theta}\|^2 \right) \right]$$

$\lambda \geq 0$ **regularization parameter:** when this is 0 we backup to OLS (no regularization at all)

$\alpha \in [0, 1]$ **tradeoff parameter:** to weight regularization penalties

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} \|\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}\|^2 + \lambda \left(\alpha \|\boldsymbol{\theta}\| + (1 - \alpha) \|\boldsymbol{\theta}\|^2 \right) \right]$$

$\lambda \geq 0$ **regularization parameter:** when this is 0 we backup to OLS (no regularization at all)

$\alpha \in [0, 1]$ **tradeoff parameter:** to weight regularization penalties

$\lambda > 0; \alpha = 1$ **Least Absolute Shrinkage and Selection Operator** or **LASSO** (L1-regularization only)

Elastic Net Framework

We consider a far more general optimization framework,
which OLS is just a special case of

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{m} ||\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}||^2 + \lambda \left(\alpha ||\boldsymbol{\theta}||_1 + (1 - \alpha) ||\boldsymbol{\theta}||^2 \right) \right]$$

$\lambda \geq 0$ **regularization parameter:** when this is 0 we backup to OLS (no regularization at all)

$\alpha \in [0, 1]$ **tradeoff parameter:** to weight regularization penalties

$\lambda > 0; \alpha = 1$ **Least Absolute Shrinkage and Selection Operator** or **LASSO** (L1-regularization only)

$\lambda > 0; \alpha = 0$ **Ridge** (L2-regularization only)

Take-Home Message of Today

- Linear Regression is a simple yet powerful tool for learning real-valued functions between feature and response variables

Take-Home Message of Today

- Linear Regression is a simple yet powerful tool for learning real-valued functions between feature and response variables
- The estimation of model's parameters is usually done via Ordinary Least Squares (OLS) by minimizing Mean Squared Error (MSE)

Take-Home Message of Today

- Linear Regression is a simple yet powerful tool for learning real-valued functions between feature and response variables
- The estimation of model's parameters is usually done via Ordinary Least Squares (OLS) by minimizing Mean Squared Error (MSE)
- OLS admits a closed-form solution which allows computing the parameters analytically via the pseudo-inverse of the feature matrix \mathbf{X}

Take-Home Message of Today


- Linear Regression is a simple yet powerful tool for learning real-valued functions between feature and response variables
- The estimation of model's parameters is usually done via Ordinary Least Squares (OLS) by minimizing Mean Squared Error (MSE)
- OLS admits a closed-form solution which allows computing the parameters analytically via the pseudo-inverse of the feature matrix \mathbf{X}
- Several quality measures: RSE, R^2 , Adjusted R^2 , etc.

Take-Home Message of Today

- Linear Regression is a simple yet powerful tool for learning real-valued functions between feature and response variables
- The estimation of model's parameters is usually done via Ordinary Least Squares (OLS) by minimizing Mean Squared Error (MSE)
- OLS admits a closed-form solution which allows computing the parameters analytically via the pseudo-inverse of the feature matrix \mathbf{X}
- Several quality measures: RSE, R^2 , Adjusted R^2 , etc.
- Regularization to prevent overfitting: Elastic Net, LASSO, Ridge

Further Readings

An Introduction to Statistical Learning [Chapter 3]



Freely available at:

[http://faculty.marshall.usc.edu/gareth-james/ISL/ISLR Seventh Printing.pdf](http://faculty.marshall.usc.edu/gareth-james/ISL/ISLR%20Seventh%20Printing.pdf)